

**Objective:**

The goal of this assignment is to create a pipeline that combines advanced OpenCV image preprocessing techniques with a deep learning model to classify images into predefined categories. You will explore both traditional image processing and modern deep learning to achieve a robust solution.

**Assignment Description****Dataset Description:**

This dataset is a curated collection of 3710 chest X-ray images, accompanied by comprehensive metadata, aimed at advancing research in medical image analysis, particularly for multi-label disease classification tasks. The dataset provides an opportunity to explore the application of machine learning in identifying and classifying multiple co-occurring chest conditions from X-ray imagery.

Dataset [Download Link]: <https://www.kaggle.com/datasets/rishabhrc/chest-x-ray-dataset>

**Tasks:****Step 1: Data Augmentation and Preprocessing (OpenCV)**

Perform the following preprocessing steps on the dataset using OpenCV:

- Resize all images to 224x224 pixels.

- Apply Gaussian noise to simulate real-world noise.

- Implement advanced augmentations like rotation, perspective transformation, and histogram equalization.

- Split the data into training, validation, and test sets.

**Step 2: Feature Extraction Using OpenCV**

Extract features from the images using advanced OpenCV techniques:

- Use ORB or SIFT to identify keypoints and descriptors.

- Save the descriptors for visualization and analysis.

**Step 3: Deep Learning Pipeline**

- Use a pre-trained convolutional neural network (e.g., ResNet50, VGG16) for classification.

- Fine-tune the model on the preprocessed images.

- Experiment with different optimizers, learning rates, and loss functions.

- Use transfer learning and freeze certain layers during training.

**Step 4: Integrate OpenCV Features with Deep Learning**

- Combine OpenCV-derived features (e.g., keypoint descriptors) with deep learning predictions to create a hybrid model.

- Evaluate whether the combination improves classification performance.

**Step 5: Evaluation**

- Evaluate the model's performance on the test set using accuracy, precision, recall, and F1-score.

- Visualize model performance using confusion matrices and ROC curves.

Compare the results of the deep learning model with and without OpenCV preprocessing.

**Deliverables:**

- A Jupyter notebook (.ipynb) with:
  - Code implementation.
  - Visualization of preprocessing steps.
  - Model architecture summary.
  - Evaluation metrics and plots.
- A report summarizing:
  - Methodology used.
  - Comparison of model performance.
  - Insights and conclusions.

**Extra Credit:**

- Implement Grad-CAM to visualize what the CNN focuses on during classification.
- Automate the entire pipeline using a Python script.

**Evaluation Criteria:**

Criteria	Weightage
Data Preprocessing & Augmentation	20%
OpenCV Feature Extraction	20%
Deep Learning Model Performance	30%
Integration of OpenCV & DL	25%
Report & Documentation	5%